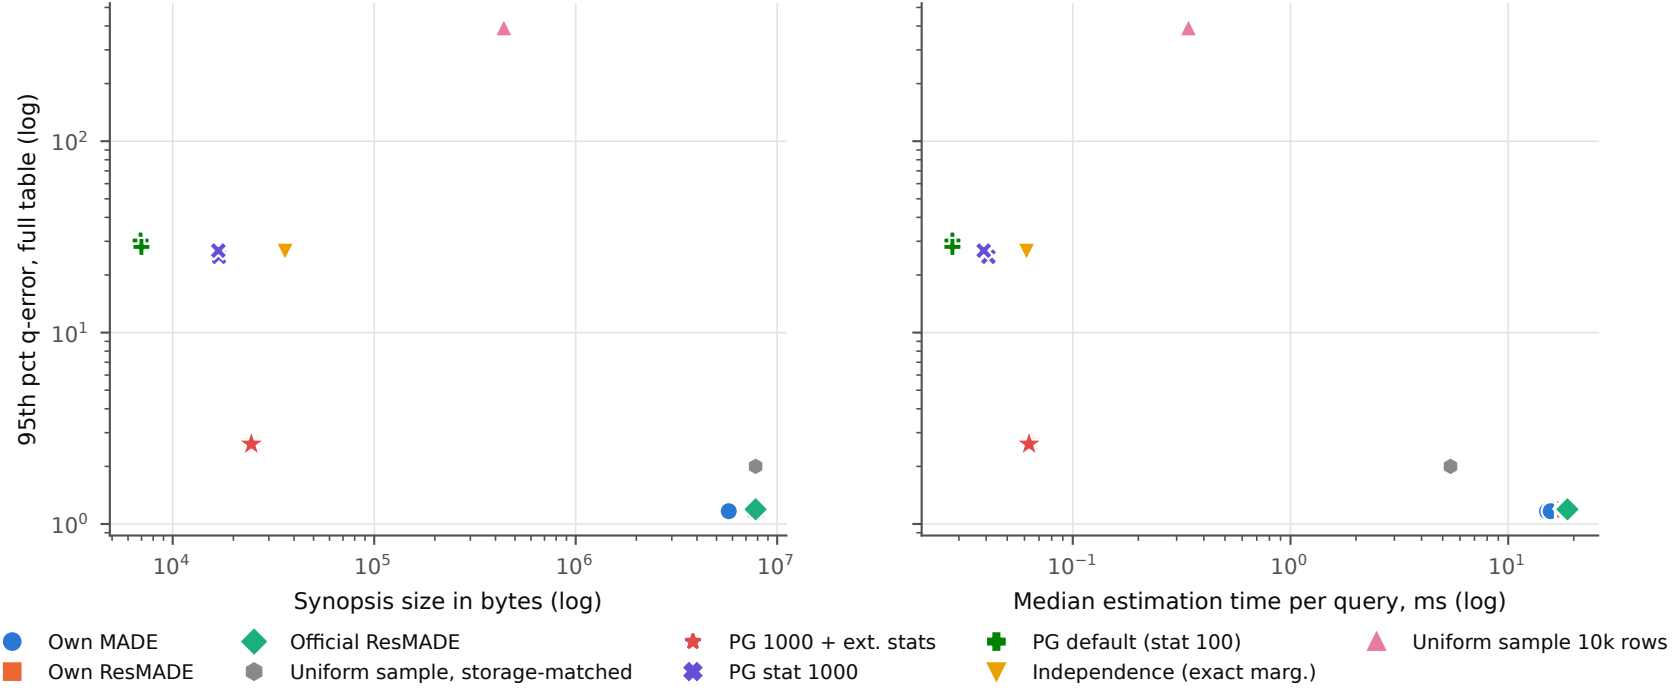

One point per training seed / ANALYZE repeat / sample seed. Size: FP32 parameters (models), int32 sample rows, PG catalog rows (pg_column_size); shared dictionaries excluded.
Time: GPU single-query inference (models), CPU (baselines), PG planner time (no execution) — different devices, not a speed-up claim.